

# Raghuraman Venkatesh

[rvenka26@uic.edu](mailto:rvenka26@uic.edu) | <https://www.linkedin.com/in/raghuraman-venkatesh-96905013a/> | 650-880-5254

## WORK EXPERIENCE

### Abbott Neuromodulation

Software Engineer II – AI R&D

Dec 2024 – Present

Dallas, Texas

- Engineered **deep learning vision algorithms** to compute 3D Human Pose Estimation of people using monocular monitored videos over time of 6 months
- Collaborated with physicians to translate clinical observations of pain-induced movement into structured 3D data frameworks for quantitative biomarker development
- Conducted comprehensive statistical analysis of pain therapy effectiveness using t-tests, correlation coefficients, and longitudinal analysis methods to identify significant patterns in patient outcomes over time
- Utilized 3D motion capture technology to map human movement and performed advanced statistical analysis (amplitude, frequency, symmetry, consistency) to characterize movement dynamics
- Deployed and managed AI models on **Google Cloud Platform (GCP)**, utilizing **Cloud Storage** for data handling and **AI Platform / Cloud Run** for on-demand and scheduled inference pipelines.

### CHIMA

AI Engineer

June 2024 – Nov 2024

San Francisco, California

- Engineered a Retrieval-Augmented Generation (RAG) System using **LangChain** and **Pinecone** to process diverse file types (CSV, Excel), enabling semantic understanding and segmentation for enhanced summaries across internal, hybrid, and external queries
- Optimized an Intelligent News Retrieval System by integrating Pinecone for semantic search and **LangChain** for orchestration, reducing runtime from 4 minutes to under 2 minutes and delivering four distinct news content types to boost engagement
- Developed an AI-powered market research system by integrating Exa Search with a LangChain-driven **RAG pipeline** and Pinecone for semantic retrieval, generating insights and visual reports from external market and internal company data

### UIC

Teaching Assistant for Machine Learning

Jan 2024 – May 2024

Chicago, Illinois

- Synchronized course content with common queries, enhancing graduate students' comprehension of intricate machine learning topics, resulting in a 17.5% increase in assignment performance

### Motorola Mobility

Computer Research Assistant

Jan 2023 – June 2024

Chicago, Illinois

- Optimized the memory footprint of the exposure fusion and low light enhancement machine learning **TFLite** model, reducing it from 2.5 GB to 1.4 GB (44% decrease)
- Innovated on the existing data-augmentation methods by creating state-of-the-art 64 augmentation blurs to mimic human hand-held blurs
- Augmentation using 24 blurs based on direction of the blurs on different exposure levels helped in reducing noise, figments and artifacts on the fused image.
- Conducted large-scale ablation study on loss functions (SSIM, GAN, MS-SSIM, Fourier, Perceptual) for image enhancement
- Engineered 34% improvement on BRISQUE and 18% on NIQE, significantly enhancing image quality obtaining state-of-the-art results

## EDUCATION

### University of Illinois (GPA 3.9/4.0)

Chicago, Illinois

Master of Science, Computer Science

Aug 2022 – May 2024

### SSN College of Engineering

Chennai, India

Bachelor of Engineering – Computer Science and Engineering (GPA 8.434 / 10.0)

Aug 2018 – June 2022

## SKILLS

**Programming:** Python, Java, C++, SQL, Scala

**ML & Data Science:** TensorFlow, PyTorch, OpenCV, Keras, NumPy, Pandas, Scikit-learn, TFLite

**Web Development:** React, Spring Boot, Flutter, Bootstrap, Tailwind, Ant-Design

**Tools & Cloud:** Docker, Kubernetes, Git, AWS, Hadoop, Django, GCP, Azure

**Generative AI & RAG:** RAG pipelines, multi-agent systems, vector databases (Pinecone, Weaviate, Chroma), prompt engineering, OpenAI API, Anthropic API, Azure OpenAI, AWS Bedrock, Exa

**Cloud & DevOps:** AWS (SageMaker, EC2, S3), Azure ,GCP (Cloud Run, Vertex AI), Docker, Kubernetes, Terraform, CI/CD

## PROJECTS

### 3D - Human Pose Estimation Dataset and comparison (OpenCV, TensorFlow, PyTorch, Python, MATLAB)

- Bridged the gap in the existing datasets by creating a new one that had more than 2 person occlusions and 10 different lighting conditions
- Increased the MPJPE score by 7.3 % to obtain better representation in both 2-D and 3-D pose estimation

### Solid Waste Classification using Deep Learning (Python, OpenCV, TensorFlow-GPU, Label-Img, Keras)

- Engineered double-stage, single-stage detectors for multi-object recognition, improving mAP and testing speed (+2 frames/sec)
- Solved the issue of lack of dataset by creating a dataset with 3200 images along with 22,500 manually annotated objects
- Exceeded benchmarks with 84/88 mAP (Faster RCNN) and 85/98 mAP (YOLOv5) using advanced Transfer Learning strategies
- Evaluated on 500 images with testing times of 0.424s (Faster RCNN) and 0.236s (YOLOv5)

## PUBLICATION

Multi Object Detection and Classification in Solid Waste Management using Region Proposal Network and YOLO model - [Gnest\\_04501](#)